

# CloudBees Agentic Software Engineer Take-Home Assessment

**Timebox:** ~1 hour

**Tools:** You are encouraged to use a coding assistant (e.g., GitHub Copilot, ChatGPT, Claude, etc.). In the interview, we'll ask you to reflect on your experience: what worked, what didn't, and where it helped or hindered.

---

## Objective

This exercise is designed to be a conversation driver in the hiring process and CloudBee's has a remote work posture so we have to drive change sometimes through written word

- Prototype **agentic workflows** with modern frameworks.
  - Translate product concepts into working demos.
  - Reason about **autonomy, reliability, and safety** in LLM-driven flows.
  - Use coding assistants as part of your workflow.
- 

## Instructions

### 1. Choose an Agentic Framework or Tool

Select one of the following (or another open-source option you prefer):

- [LangChain](#)
- [CrewAI](#)
- [Agnor](#)
- MCP-enabled environments (e.g., Goose, Claude Desktop)

### 2. Task: Prototype a Simple Agentic Flow

Build a small demo that shows how an **AI agent can autonomously interact with tools or APIs**. For example, pick one of these scenarios (or define your own):


- An agent that queries a public API (e.g., weather, GitHub issues) and summarizes results for a developer.
- A multi-step flow where an agent retrieves documentation, reasons over it, and produces a suggested action (e.g., config update, deployment step).

- A small “team of agents” where one agent retrieves information and another validates or reformats it before output.
  - 3. Keep it lightweight but demonstrate autonomy, reasoning, and tool usage.
  - 4. **Deliverables**  
Please provide:
    - **Code:** Your implementation (script, notebook, or small project).
    - **README.md:** Short explanation (1–2 paragraphs) describing:
      - What workflow you built and why.
      - How to run it.
      - Any assumptions or limitations.
    - **Reflection (5–10 sentences):** Share your experience using a coding assistant. Specifically:
      - Did it meaningfully help design the agent flow?
      - Did it generate incorrect or unsafe suggestions?
      - Where was it most/least useful?
- 

## Submission

Share a link(<https://github.com/jburtbee>) to your repo/branch and send it prior to the next interview session.

---

 **Tip:** Focus on creativity and clarity — we’re not looking for production systems, just a well-scoped demonstration of agentic reasoning in action.